

# Routerlab

Summer semester 2018

Worksheet 3  
Group 08

Valentin Franck, 364066  
Nikhil Singh, 387694

Pages: 23

Submission Date: May 28, 2018

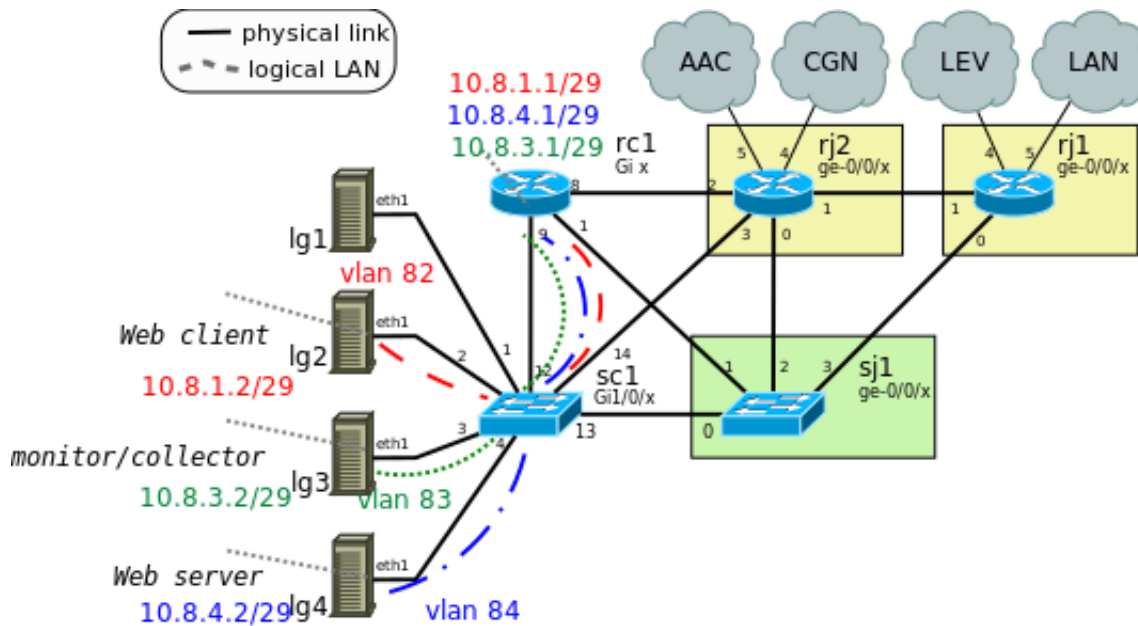


Figure 1: The IPv4 topology.

## Question 1

1a

See Figure 1.

1b

We also assigned IPv6 addresses and provide the traceroute from web client to web server:

```
root@group08-lg2:~# traceroute 10.8.4.2
traceroute to 10.8.4.2 (10.8.4.2), 30 hops max, 60 byte packets
 1  10.8.1.1 (10.8.1.1)  0.772 ms  0.938 ms  0.905 ms
 2  10.8.4.2 (10.8.4.2)  1.162 ms  1.133 ms  1.089 ms

root@group08-lg2:~# traceroute fc00:470:525b:f804::2
traceroute to fc00:470:525b:f804::2 (fc00:470:525b:f804::2), 30 hops
max, 80 byte packets
 1  fc00:470:525b:f801::1 (fc00:470:525b:f801::1)  0.841 ms  0.784 ms
    0.770 ms
 2  fc00:470:525b:f804::2 (fc00:470:525b:f804::2)  1.289 ms  1.243 ms
    1.201 ms
```

The configuration:

SC1:

```
interface GigabitEthernet1/0/2
description lg2
switchport access vlan 82
switchport mode access
!
```

```
interface GigabitEthernet1/0/3
description lg3
```

```

switchport access vlan 83
switchport mode access
!

interface GigabitEthernet1/0/4
description lg4
switchport access vlan 84
switchport mode access
!

interface GigabitEthernet1/0/12
description lev-rc1
switchport trunk allowed vlan 82-84
switchport mode trunk
!

RC1:

interface GigabitEthernet9
description lev-scl
no ip address
duplex auto
speed auto
!
interface GigabitEthernet9.82
encapsulation dot1Q 82
ip address 10.8.1.1 255.255.255.248
ipv6 address FC00:470:525B:F801::1/64
no cdp enable
!
interface GigabitEthernet9.83
encapsulation dot1Q 83
ip address 10.8.3.1 255.255.255.248
ipv6 address FC00:470:525B:F803::1/64
no cdp enable
!
interface GigabitEthernet9.84
encapsulation dot1Q 84
ip address 10.8.4.1 255.255.255.248
ipv6 address FC00:470:525B:F804::1/64
no cdp enable
!

LG2:

root@group08-lg2:~# ip a s
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
    group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
    state UP group default qlen 1000
    link/ether 00:16:3e:af:08:20 brd ff:ff:ff:ff:ff:ff
    inet 172.16.0.209/20 brd 172.16.15.255 scope global eth0
        valid_lft forever preferred_lft forever

```

```

        inet6 fe80::216:3eff:feaf:820/64 scope link
        valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
   state UP group default qlen 1000
   link/ether 00:16:3e:af:08:21 brd ff:ff:ff:ff:ff:ff
   inet 10.8.1.2/29 scope global eth1
       valid_lft forever preferred_lft forever
   inet6 fc00:470:525b:f801:216:3eff:feaf:821/64 scope global
       mngtmpaddr dynamic
       valid_lft 2491291sec preferred_lft 504091sec
   inet6 fc00:470:525b:f801::2/64 scope global
       valid_lft forever preferred_lft forever
   inet6 fe80::216:3eff:feaf:821/64 scope link
       valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group
   default qlen 1000
   link/ether 00:16:3e:af:08:22 brd ff:ff:ff:ff:ff:ff

```

```
root@group08-lg2:~# route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use
Iface						
10.8.0.0	10.8.1.1	255.255.255.248	UG	0	0	0
eth1						
10.8.1.0	0.0.0.0	255.255.255.248	U	0	0	0
eth1						
10.8.2.0	10.8.1.1	255.255.255.248	UG	0	0	0
eth1						
10.8.3.0	10.8.1.1	255.255.255.248	UG	0	0	0
eth1						
10.8.4.0	10.8.1.1	255.255.255.248	UG	0	0	0
eth1						
172.16.0.0	0.0.0.0	255.255.240.0	U	0	0	0
eth0						
172.16.255.0	lion.routerlab	255.255.255.0	UG	0	0	0
eth0						

```
LG3
```

```
root@group08-lg3:~# ip a s
```

```

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
   group default
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
   inet 127.0.0.1/8 scope host lo
       valid_lft forever preferred_lft forever
   inet6 ::1/128 scope host
       valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
   state UP group default qlen 1000
   link/ether 00:16:3e:af:08:30 brd ff:ff:ff:ff:ff:ff
   inet 172.16.0.210/20 brd 172.16.15.255 scope global eth0
       valid_lft forever preferred_lft forever
   inet6 fe80::216:3eff:feaf:830/64 scope link
       valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
   state UP group default qlen 1000
   link/ether 00:16:3e:af:08:31 brd ff:ff:ff:ff:ff:ff

```

```

    inet 10.8.3.2/29 scope global eth1
        valid_lft forever preferred_lft forever
    inet6 fc00:470:525b:f803::2/64 scope global
        valid_lft forever preferred_lft forever
    inet6 fc00:470:525b:f803:216:3eff:feaf:831/64 scope global
        mngtmpaddr dynamic
        valid_lft 2591957sec preferred_lft 604757sec
4: eth2: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group
    default qlen 1000
    link/ether 00:16:3e:af:08:32 brd ff:ff:ff:ff:ff:ff

```

```
root@group08-lg3:~# route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use
Iface						
10.8.1.0	10.8.3.1	255.255.255.248	UG	0	0	0
eth1						
10.8.2.0	10.8.3.1	255.255.255.248	UG	0	0	0
eth1						
10.8.3.0	0.0.0.0	255.255.255.248	U	0	0	0
eth1						
10.8.4.0	10.8.3.1	255.255.255.248	UG	0	0	0
eth1						
172.16.0.0	0.0.0.0	255.255.240.0	U	0	0	0
eth0						
172.16.255.0	lion.routerlab	255.255.255.0	UG	0	0	0
eth0						

```
LG4
```

```
root@group08-lg4:~# ip a s
```

```

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
    group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
    state UP group default qlen 1000
    link/ether 00:16:3e:af:08:40 brd ff:ff:ff:ff:ff:ff
    inet 172.16.0.211/20 brd 172.16.15.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::216:3eff:feaf:840/64 scope link
        valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast
    state UP group default qlen 1000
    link/ether 00:16:3e:af:08:41 brd ff:ff:ff:ff:ff:ff
    inet 10.8.4.2/29 scope global eth1
        valid_lft forever preferred_lft forever
    inet6 fc00:470:525b:f804::2/64 scope global
        valid_lft forever preferred_lft forever
    inet6 fe80::216:3eff:feaf:841/64 scope link
        valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group
    default qlen 1000
    link/ether 00:16:3e:af:08:42 brd ff:ff:ff:ff:ff:ff

```

TCP Bandwidth	MSS(Client)	MSS(Server)	TCP window size(Client)	TCP window size(Server)
22.1 Gbit/s	65483 bytes	21888 bytes	2.50 MByte	85.3 KByte
22.0 Gbit/s	65483 bytes	21888 bytes	2.50 MByte	85.3 KByte
21.8 Gbit/s	65483 bytes	21888 bytes	2.50 MByte	85.3 KByte
22.1 Gbit/s	65483 bytes	21888 bytes	2.50 MByte	85.3 KByte
22.5 Gbit/s	65483 bytes	21888 bytes	2.50 MByte	85.3 KByte

Table 1: Iperf server and client running on the same machine.

```
root@group08-lg4:~# route
```

```
Kernel IP routing table
```

Destination Iface	Gateway	Genmask	Flags	Metric	Ref	Use
10.8.1.0 eth1	10.8.4.1	255.255.255.248	UG	0	0	0
10.8.2.0 eth1	10.8.4.1	255.255.255.248	UG	0	0	0
10.8.3.0 eth1	10.8.4.1	255.255.255.248	UG	0	0	0
10.8.4.0 eth1	0.0.0.0	255.255.255.248	U	0	0	0
172.16.0.0 eth0	0.0.0.0	255.255.240.0	U	0	0	0
172.16.255.0 eth0	lion.routerlab	255.255.255.0	UG	0	0	0

## Question 2

### 2a

We used screen to launch the iperf server and have it run in background while launching the client. Here is the output of the server and the client, used to create Table 1.

```
root@group08-lg2:~# iperf -s -m
```

```
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)
```

```
[ 4] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 57121
[ ID] Interval      Transfer      Bandwidth
[ 4] 0.0-10.0 sec  25.8 GBytes  22.1 Gbits/sec
[ 4] MSS size 21888 bytes (MTU 21928 bytes, unknown interface)
[ 5] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 57122
[ 5] 0.0-10.0 sec  25.6 GBytes  22.0 Gbits/sec
[ 5] MSS size 21888 bytes (MTU 21928 bytes, unknown interface)
[ 4] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 57123
[ 4] 0.0-10.0 sec  25.3 GBytes  21.8 Gbits/sec
[ 4] MSS size 21888 bytes (MTU 21928 bytes, unknown interface)
[ 5] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 57124
[ 5] 0.0-10.0 sec  25.8 GBytes  22.1 Gbits/sec
[ 5] MSS size 21888 bytes (MTU 21928 bytes, unknown interface)
[ 4] local 127.0.0.1 port 5001 connected with 127.0.0.1 port 57125
[ 4] 0.0-10.0 sec  26.1 GBytes  22.4 Gbits/sec
[ 4] MSS size 21888 bytes (MTU 21928 bytes, unknown interface)
```

```
root@group08-lg2:~# iperf -c localhost -m
```

---

```
Client connecting to localhost, TCP port 5001
TCP window size: 2.50 MByte (default)
```

---

```
[ 3] local 127.0.0.1 port 57121 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.0 sec  25.8 GBytes  22.1 Gbits/sec
[ 3] MSS size 65483 bytes (MTU 65523 bytes, unknown interface)
root@group08-lg2:~# iperf -c localhost -m
```

---

```
Client connecting to localhost, TCP port 5001
TCP window size: 2.50 MByte (default)
```

---

```
[ 3] local 127.0.0.1 port 57122 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.0 sec  25.6 GBytes  22.0 Gbits/sec
[ 3] MSS size 65483 bytes (MTU 65523 bytes, unknown interface)
root@group08-lg2:~# iperf -c localhost -m
```

---

```
Client connecting to localhost, TCP port 5001
TCP window size: 2.50 MByte (default)
```

---

```
[ 3] local 127.0.0.1 port 57123 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.0 sec  25.3 GBytes  21.8 Gbits/sec
[ 3] MSS size 65483 bytes (MTU 65523 bytes, unknown interface)
root@group08-lg2:~# iperf -c localhost -m
```

---

```
Client connecting to localhost, TCP port 5001
TCP window size: 2.50 MByte (default)
```

---

```
[ 3] local 127.0.0.1 port 57124 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.0 sec  25.8 GBytes  22.1 Gbits/sec
[ 3] MSS size 65483 bytes (MTU 65523 bytes, unknown interface)
root@group08-lg2:~# iperf -c localhost -m
```

---

```
Client connecting to localhost, TCP port 5001
TCP window size: 2.50 MByte (default)
```

---

```
[ 3] local 127.0.0.1 port 57125 connected with 127.0.0.1 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.0 sec  26.1 GBytes  22.5 Gbits/sec
[ 3] MSS size 65483 bytes (MTU 65523 bytes, unknown interface)
```

## 2b

Here is the output of the server and the client, used to create Table 2.

```
root@group08-lg4:~# iperf -s -m
```

---

```
Server listening on TCP port 5001
TCP window size: 85.3 KByte (default)
```

---

```
[ 4] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44988
```

TCP Bandwidth (C.)	Bandwidth (S.)	MSS	TCP window size (C.)	window size (S.)
706 Mbits/sec	705 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte
642 Mbits/sec	642 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte
896 Mbits/sec	894 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte
676 Mbits/sec	676 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte
778 Mbits/sec	776 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte
626 Mbits/sec	625 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte
633 Mbits/sec	632 Mbits/sec	1448 bytes	85.0 KByte	85.3 KByte

Table 2: Iperf server and client running on LG2 and LG4 respectively.

```
[ ID] Interval      Transfer      Bandwidth
[ 4]  0.0-10.0 sec  842 MBytes  705 Mbits/sec
[ 4] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
[ 5] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44989
[ 5]  0.0-10.0 sec  766 MBytes  642 Mbits/sec
[ 5] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
[ 4] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44990
[ 4]  0.0-10.0 sec  1.04 GBytes  894 Mbits/sec
[ 4] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
[ 5] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44991
[ 5]  0.0-10.0 sec  806 MBytes  676 Mbits/sec
[ 5] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
[ 4] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44992
[ 4]  0.0-10.0 sec  928 MBytes  776 Mbits/sec
[ 4] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
[ 5] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44993
[ 5]  0.0-10.0 sec  746 MBytes  625 Mbits/sec
[ 5] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
[ 4] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44994
[ 4]  0.0-10.0 sec  754 MBytes  632 Mbits/sec
[ 4] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
```

```
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

```
Client connecting to 10.8.4.2, TCP port 5001
TCP window size: 85.0 KByte (default)
```

```
[ 3] local 10.8.1.2 port 44988 connected with 10.8.4.2 port 5001
```

```
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec  842 MBytes  706 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
```

```
root@group08-lg2:~#
```

```
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

```
Client connecting to 10.8.4.2, TCP port 5001
TCP window size: 85.0 KByte (default)
```

```
[ 3] local 10.8.1.2 port 44989 connected with 10.8.4.2 port 5001
```

```
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec  766 MBytes  642 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
```

```
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

```
Client connecting to 10.8.4.2, TCP port 5001
```



TCP window size: 85.0 KByte (default)

---

```
[ 3] local 10.8.1.2 port 44990 connected with 10.8.4.2 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec  1.04 GBytes   896 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

---

Client connecting to 10.8.4.2, TCP port 5001  
 TCP window size: 85.0 KByte (default)

---

```
[ 3] local 10.8.1.2 port 44991 connected with 10.8.4.2 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec   806 MBytes   676 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

---

Client connecting to 10.8.4.2, TCP port 5001  
 TCP window size: 85.0 KByte (default)

---

```
[ 3] local 10.8.1.2 port 44992 connected with 10.8.4.2 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec   928 MBytes   778 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

---

Client connecting to 10.8.4.2, TCP port 5001  
 TCP window size: 85.0 KByte (default)

---

```
[ 3] local 10.8.1.2 port 44993 connected with 10.8.4.2 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec   746 MBytes   626 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
root@group08-lg2:~# iperf -c 10.8.4.2 -m
```

---

Client connecting to 10.8.4.2, TCP port 5001  
 TCP window size: 85.0 KByte (default)

---

```
[ 3] local 10.8.1.2 port 44994 connected with 10.8.4.2 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec   754 MBytes   633 Mbits/sec
[ 3] MSS size 1448 bytes (MTU 1500 bytes, ethernet)
```

---

## 2c

Yes we observe the differences in the values from Q2(a) and (b). The parameters that have changed are TCP Bandwidth, MSS and TCP Window size. TCP Bandwidth in Q2(a) was the same in both client and server side which is way higher than Q2(b) where TCP Bandwidth is minutely less in some cases on the server side than the client side. We observed different MSS on client and server side in Q2(a). On the client side it is approximately three times more than the MSS of server which is also more than 10 times the MSS in Q2(b). TCP Window size on the server is by default same on both Q2(a) and Q2(b). However, TCP window size of the client side in Q2(a) is 2.50 MByte which is quite larger than Q2(b) which is 85.0 KByte.

The reason for these changes is that in the latter case the data actually has to be transmitted over the network, causing the typical delays (transmission, propagation, processing and queuing delay). In the first case it was just sent to the loopback interface, which is a virtual interface

emulated by the kernel and therefore does not have the physical limitations of a real network.

## 2d

Nagle's Algorithm is used to reduce the overhead in TCP by merging various smaller packets into one larger packet. This algorithm might not be suitable for interactive real time applications, e.g. multiplayer games or application layer protocols like Telnet, which often only expects a single byte (one key stroke). If this was delayed by Nagle's algorithm, it would affect the user experience and really slow down the protocol. So if one wants to simulate such environments, where it makes sense to immediately send even small packets, one needs to disable Nagle's algorithm.

Bidirectional measurements make most sense, when data really is sent in both directions. The current client-server architecture of many distributed systems results in very asymmetric traffic. So the client usually sends very little data and the server a lot of data (consider multimedia streaming). There might even be scenarios, in which data is only sent in one direction (using a connectionless protocol like UDP. An example might be a network with a lot of real time sensor data, that is sent via UDP (fire and forget). In such a network it does not makes sense to perform bidirectional bandwidth tests.

The output from the bidirectional measurement of the client:

```
root@group08-lg2:~# iperf -c 10.8.4.2 -d
```

---

```
Server listening on TCP port 5001
TCP window size: 85.3 KByte (default)
```

---



---

```
Client connecting to 10.8.4.2, TCP port 5001
TCP window size: 264 KByte (default)
```

---

```
[ 3] local 10.8.1.2 port 44913 connected with 10.8.4.2 port 5001
[ 5] local 10.8.1.2 port 5001 connected with 10.8.4.2 port 34324
[ ID] Interval      Transfer      Bandwidth
[ 3]  0.0-10.0 sec   325 MBytes   272 Mbits/sec
[ 5]  0.0-10.1 sec   328 MBytes   274 Mbits/sec
```

The output from the server:

```
root@group08-lg4:~# iperf -s
```

---

```
Server listening on TCP port 5001
TCP window size: 85.3 KByte (default)
```

---



---

```
[ 4] local 10.8.4.2 port 5001 connected with 10.8.1.2 port 44913
```

---



---

```
Client connecting to 10.8.1.2, TCP port 5001
TCP window size: 162 KByte (default)
```

---

```
[ 6] local 10.8.4.2 port 34324 connected with 10.8.1.2 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 6]  0.0-10.0 sec   328 MBytes   275 Mbits/sec
[ 4]  0.0-10.1 sec   325 MBytes   271 Mbits/sec
```

## Question 3

### 3a

As per the Harpoon paper published by Joel Sommers, Hyungsuk Kim and Paul Barford tools like iperf lack the richness and diversity of packets streams observed in the live internet. Harpoon

targets traffic generation at the IP flow level. It gives router designers and network operators insight of system demeanor under real operating conditions.

### 3b

Harpoon combines seven distributional models: file size, inter-file request time, session duration, inter-session request time, IP range, user ON time and number of active users. These model work at different levels in the hierarchy: user, session and file level. These parameters are drawn from different distributions.

- file size - the size of the file that is transferred
- inter-file request time - the time between file consecutive file transfer requests.
- IP spatial distribution - the space from which IP addresses of destination and source are drawn for a session
- inter-session start times - the time interval that separates the start of consecutive sessions
- session duration - the time interval during which file transfer requests between source and destination take place
- user ON time - the time a user is active (drawn from a distribution)
- active users - the number of active users, which may vary over time as it does in the Internet

In order to simulate UDP traffic harpoon simply transfers datagram of a fixed size at a constant bit rate, while TCP file transfers are controlled by protocol dynamics (run at end hosts).

### 3c

We explained these parameters in Question 3b: inter-connection time is the same as inter-file request time and active sessions corresponds to the number of active users. There have been some changes between different versions of harpoon, which explain the differences in architecture and names. In the terminology used in our config files harpoon only consists of a two level architecture.

### 3d

```
root@group08-lg4:/usr/local/harpoon/run_harpoon.sh -f /root/profiles/
web-server.xml -v 10 -w 300
```

```
root@group08-lg4:/usr/local/harpoon# netstat -an
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 0.0.0.0:8180             0.0.0.0:*               LISTEN
tcp        0      0 0.0.0.0:22               0.0.0.0:*               LISTEN
tcp        0      0 0.0.0.0:10000            0.0.0.0:*               LISTEN
tcp6       0      0 :::22                   :::*                    LISTEN
udp        0      0 0.0.0.0:68              0.0.0.0:*
Active UNIX domain sockets (servers and established)
Proto RefCnt Flags       Type       State       I-Node     Path
unix   3      [ ]       DGRAM     7752       /run/systemd
/notify
```

```

unix  2      [ ]          DGRAM          7754      /run/systemd
/cgroups-agent
unix  2      [ ACC ]      STREAM          LISTENING    7758      /run/systemd
/private
unix  2      [ ]          DGRAM          7765      /run/systemd
/journal/syslog
unix  4      [ ]          DGRAM          7776      /run/systemd
/journal/dev-log
unix  2      [ ACC ]      STREAM          LISTENING    7779      /run/systemd
/fsck.progress
unix  2      [ ACC ]      SEQPACKET      LISTENING    8042      /run/udev/
control
unix  2      [ ACC ]      STREAM          LISTENING    7789      /run/systemd
/journal/stdout
unix  5      [ ]          DGRAM          7791      /run/systemd
/journal/socket
unix  3      [ ]          DGRAM          8878
unix  2      [ ]          DGRAM          240184
unix  3      [ ]          STREAM          CONNECTED    9344
unix  2      [ ]          DGRAM          8593
unix  3      [ ]          STREAM          CONNECTED    11029     /run/systemd
/journal/stdout
unix  3      [ ]          DGRAM          7757
unix  3      [ ]          DGRAM          8681
unix  3      [ ]          STREAM          CONNECTED    8492
unix  2      [ ]          DGRAM          137863
unix  2      [ ]          DGRAM          8860
unix  3      [ ]          STREAM          CONNECTED    11028
unix  2      [ ]          DGRAM          9471
unix  3      [ ]          DGRAM          8879
unix  2      [ ]          DGRAM          8047
unix  3      [ ]          DGRAM          8680
unix  3      [ ]          DGRAM          7756
unix  2      [ ]          DGRAM          9436
unix  3      [ ]          STREAM          CONNECTED    8677
unix  3      [ ]          DGRAM          8880
unix  3      [ ]          STREAM          CONNECTED    9345      /run/systemd
/journal/stdout
unix  3      [ ]          STREAM          CONNECTED    8678      /run/systemd
/journal/stdout
unix  3      [ ]          STREAM          CONNECTED    8494      /run/systemd
/journal/stdout
unix  3      [ ]          DGRAM          8881

```

```

root@group08-lg2:/usr/local/harpoon/run_harpoon.sh -f /root/profiles/
web-client.xml -v 10 -w 300 -c

```

As we can see in the screenshot in Figure 2 generated on the harpoon client 22130 kbps of incoming traffic and 237 kbp of outgoing traffic is generated. Of course these numbers vary slightly with the server sending about a hundred times more traffic than the client.

### 3e

```

root@group08-lg2:/usr/local/harpoon# iptraf-ng

```

We show the detailed interface statistics in Figure 2.

```

lev-scl#show interfaces Gi1/0/12 | include bits
30 second input rate 38167000 bits/sec, 3918 packets/sec

```

```

iptraf-ng 1.1.4
Statistics for eth1

```

	Total Packets	Total Bytes	Incoming Packets	Incoming Bytes	Outgoing Packets	Outgoing Bytes
Total:	42034	165691k	20593	164564k	21441	1126914
IPv4:	42034	165691k	20593	164564k	21441	1126914
IPv6:	0	0	0	0	0	0
TCP:	42034	165691k	20593	164564k	21441	1126914
UDP:	0	0	0	0	0	0
ICMP:	0	0	0	0	0	0
Other IP:	0	0	0	0	0	0
Non-IP:	0	0	0	0	0	0
Total rates:	22368.22 kbps			Broadcast packets:		0
	1059 pps			Broadcast bytes:		0
Incoming rates:	22130.80 kbps					
	497 pps					
Outgoing rates:	237.41 kbps			IP checksum errors:		0
Elapsed time:	0	561 pps				

X-exit

Figure 2: The iptraf output showing detailed interface statistics.

```

30 second output rate 38147000 bits/sec , 3917 packets/sec
lev-scl#show interfaces Gi1/0/2 | include bits
30 second input rate 317000 bits/sec , 562 packets/sec
30 second output rate 31834000 bits/sec , 2723 packets/sec
lev-scl#show interfaces Gi1/0/3 | include bits
30 second input rate 0 bits/sec , 0 packets/sec
30 second output rate 11000 bits/sec , 1 packets/sec
lev-scl#show interfaces Gi1/0/4 | include bits
30 second input rate 26264000 bits/sec , 2266 packets/sec
30 second output rate 302000 bits/sec , 537 packets/sec

```

```

lev-rc1#show interfaces Gi9 | include bits
30 second input rate 19688000 bits/sec , 2167 packets/sec
30 second output rate 19702000 bits/sec , 2168 packets/sec

```

### 3f

We simply changed the number of active sessions trying different values and using the python tool to figure out different values for different target traffic rates:

```

root@group08-lg2:~# python harpoon-master/selfconf/harpoon_reconf.py -d
-c profiles/web-client.xml -s profiles/web-server.xml -i 300 -r
2000000
target volume: 1500000000.0
interval duration: 300
client conf file: profiles/web-client.xml
server conf file: profiles/web-server.xml
targetbytes 1500000000.0 simbytes 206930175 median 2 mean 1 stdev
0.519809599952 max 3 flows 2146
number of sessions should be 1 to achieve volume of 1500000000 bytes
(2000000.0 bits/sec)

```

We used the following numbers of sessions. We left the remaining parameters untouched. Our

Number of active sessions	Total (kbps)	Incoming (kbps)	Outgoing (kbps)
1	2057	2039	18
100	175426	173849	1577
500	44618	43887	731
1000	121360	120138	1222
5000	85882	84901	980

Table 3: Different numbers of active sessions and the reported data rates (by iptraf-ng at the client).

data rates are shown in Table 3.

## Question 4

### 4a

```

lev-rc1(config-if)#ip address 100.100.100.100 255.255.255.255
lev-rc1(config)#interface Loopback 0
lev-rc1(config-if)#no shutdown
lev-rc1(config)#ip flow-export source Loopback0
lev-rc1(config)#ip flow-export destination 10.8.3.2 7777
lev-rc1(config)#interface Gi9.82
lev-rc1(config-subif)#ip flow ingress

lev-rc1#sh ip flow export
Flow export v1 is enabled for main cache
  Export source and destination details :
    VRF ID : Default
      Source(1)      100.100.100.100 (Loopback0)
      Destination(1) 10.8.3.2 (7777)
  Version 1 flow records
    118876 flows exported in 4962 udp datagrams
    0 flows failed due to lack of export packet
    0 export packets were sent up to process level
    0 export packets were dropped due to no fib
    0 export packets were dropped due to adjacency issues
    0 export packets were dropped due to fragmentation failures
    0 export packets were dropped due to encapsulation fixup failures

```

### 4b

Note, that we are listening to all flows. We could restrict it to flow coming from 100.100.100.100, though.

```

root@group08-lg3:~# tail /etc/flow-tools/flow-capture.conf
# Store flows in /var/flow/mysecondrouter. Rotate files every
# 5 minutes.
# -w /var/flow/mysecondrouter -n 275 0/10.3.2.6/3002

# Example 3:
# Same as above, but only listen at address 10.3.2.5, and store
# files under 'YYYY/YYYY-MM/YYYY-MM-DD' directories.
# -w /var/flow/mysecondrouter -n 275 -N 3 10.3.2.5/10.3.2.6/3002

-w /tmp/flows 0/0/7777

```

```
root@group08-lg3:~# /etc/init.d/flow-capture start
```

```
root@group08-lg3:~# flow-cat < /tmp/flows/2018/2018-05/2018-05-27/ft-
v01.2018-05-27.105055+0000 | flow-report
```

```
#  ———— Report Information ————
# build-version:      flow-tools 0.68
# name:               default
# type:               summary-detail
# options:            +header,+xheader,+totals
# fields:             +other
# records:            0
# first-flow:         1527416860 Sun May 27 10:27:40 2018
# last-flow:          1527417338 Sun May 27 10:35:38 2018
# now:                1527418823 Sun May 27 11:00:23 2018
#
# mode:               normal
# capture hostname:   group08-lg3
# capture start:      Sun May 27 10:50:55 2018
# capture end:        Sun May 27 11:00:00 2018
# capture period:     545 seconds
# compress:           on
# byte order:         little
# stream version:     3
# export version:     1
# lost flows:         0
# corrupt packets:    0
# sequencer resets:   0
# capture flows:      20040
#
# [ '/usr/bin/flow-rptfmt ', '-f', 'ascii ' ]
Ignores:              0
Total Flows:           20040
Total Octets:          19564567
Total Packets:         371887
Total Duration (ms):   166928
Real Time:             1527417338
Average Flow Time:     8.000000
Average Packets/Second: 52.000000
Average Flows/Second:  976.000000
Average Packets/Flow:  18.000000
Flows/Second:          0.115350
Flows/Second (real):   0.000013
```

Average IP packet size distribution:

1-32	64	96	128	160	192	224	256	288	320	352	384	416
448	480											
.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000
.000	.000											
512	544	576	1024	1536	2048	2560	3072	3584	4096	4608		
.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000		

Packets per flow distribution:

1	2	4	8	12	16	20	24	28	32	36	40	44
	48	52										
.000	.000	.000	.287	.407	.117	.060	.031	.020	.013	.008	.007	.006
.005	.005											
60	100	200	300	400	500	600	700	800	900	>900		
.007	.015	.007	.002	.001	.000	.000	.000	.000	.000	.001		

Octets per flow distribution:

32	64	128	256	512	1280	2048	2816	3584	4352	5120	5888	6656
7424	8192											
.000	.000	.000	.000	.411	.491	.047	.018	.011	.006	.004	.002	.001
.001	.001											
8960	9728	10496	11264	12032	12800	13568	14336	15104	15872	>15872		
.001	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.002	

Flow Time Distribution (ms):

10	50	100	200	500	1000	2000	3000	4000	5000	6000	7000	8000
9000	10000											
.892	.098	.005	.001	.001	.001	.001	.000	.000	.000	.000	.000	.000
.000	.000											
12000	14000	16000	18000	20000	22000	24000	26000	28000	30000	>30000		
.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	.000	

In the flow-report output we can see some information regarding the flow such as when it was captured, how many flows, packets, octets etc. were captured, which version was used and so on. Most importantly we can see how many, packets were captured, what was the packet size, how many bytes per flow were captured and how long lasted every flow. We can also see the distribution of these values apart from the absolute numbers and the averages.

#### 4c

In Figure 3, 4 and 5 we show the graphs for the standard configuration.

In Figure 6, 7 and 8 the graphs for the maximum throughput are shown. Unfortunately there seems to be a delay of about 30 minutes until the flows captured at the router actually are sent to the monitor. Therefore and due to our ending reservation we may not have captured a complete round of 5 minutes of maximum throughput. We provide our graphs anyway with what we have got.

```

root@group08-lg3:/tmp/flows/2018/2018-05/2018-05-28# flow-cat < tmp-v01
.2018-05-28.174622+0000 | flow-report
# ----- Report Information -----
# build-version:      flow-tools 0.68
# name:               default
# type:               summary-detail
# options:             +header,+xheader,+totals
# fields:             +other
# records:            0
# first-flow:         1527528186 Mon May 28 17:23:06 2018
# last-flow:          1527528598 Mon May 28 17:29:58 2018
# now:                1527530005 Mon May 28 17:53:25 2018
#
# mode:               streaming

```



```
# compress:                off
# byte order:              little
# stream version:          3
# export version:          1
#
# [ '/usr/bin/flow-rptfmt', '-f', 'ascii' ]
Ignores:                   0
Total Flows:               71899
Total Octets:              66190891
Total Packets:             1240568
Total Duration (ms):       12457552
Real Time:                 1527528598
Average Flow Time:         173.000000
Average Packets/Second:    53.000000
Average Flows/Second:      920.000000
Average Packets/Flow:      17.000000
Flows/Second:              5.239928
Flows/Second (real):       0.000047
```

Average IP packet size distribution:

```
1-32   64   96  128  160  192  224  256  288  320  352  384  416
    448  480
.000 .000 .000 .000 .000 .000 .000 .000 .000 .000 .000 .000 .000
.000 .000
```

```
512  544  576 1024 1536 2048 2560 3072 3584 4096 4608
.000 .000 .000 .000 .000 .000 .000 .000 .000 .000 .000
```

Packets per flow distribution:

```
1    2    4    8    12   16   20   24   28   32   36   40   44
    48   52
.000 .000 .000 .450 .278 .094 .052 .032 .020 .013 .010 .007 .006
.004 .004
```

```
60  100  200  300  400  500  600  700  800  900 >900
.006 .012 .007 .002 .001 .001 .000 .000 .000 .000 .001
```

Octets per flow distribution:

```
32   64  128  256  512 1280 2048 2816 3584 4352 5120 5888 6656
    7424 8192
.000 .000 .000 .000 .565 .337 .050 .017 .009 .005 .003 .002 .001
.001 .001
```

```
8960 9728 10496 11264 12032 12800 13568 14336 15104 15872 >15872
.001 .001 .001 .000 .000 .000 .000 .000 .000 .000 .003
```

Flow Time Distribution (ms):

```
10    50  100  200  500 1000 2000 3000 4000 5000 6000 7000 8000
    9000 10000
.209 .384 .006 .138 .208 .029 .020 .003 .002 .001 .000 .001 .000
.000 .000
```

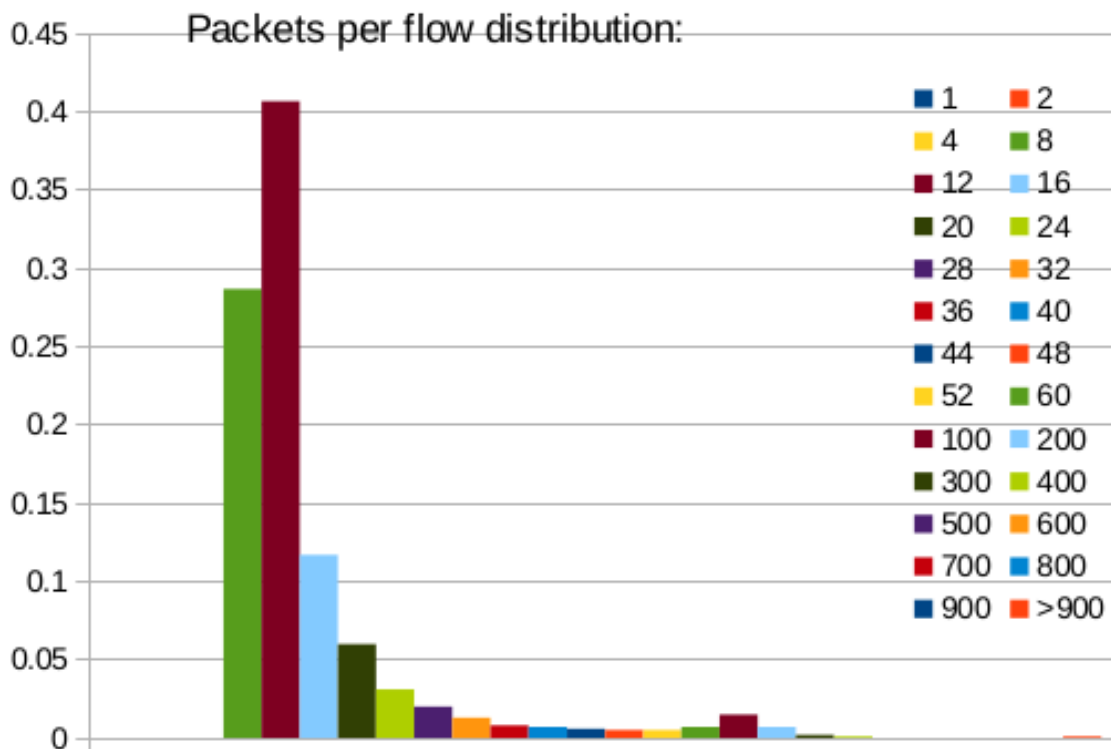


Figure 3: The packets per flow distribution with standard configuration.

```
12000 14000 16000 18000 20000 22000 24000 26000 28000 30000 >30000
.000 .000 .000 .000 .000 .000 .000 .000 .000 .000 .000
```

## Question 5

### 5a

```
lev-rc1(config)#snmp-server community OUR_SECRET RO
```

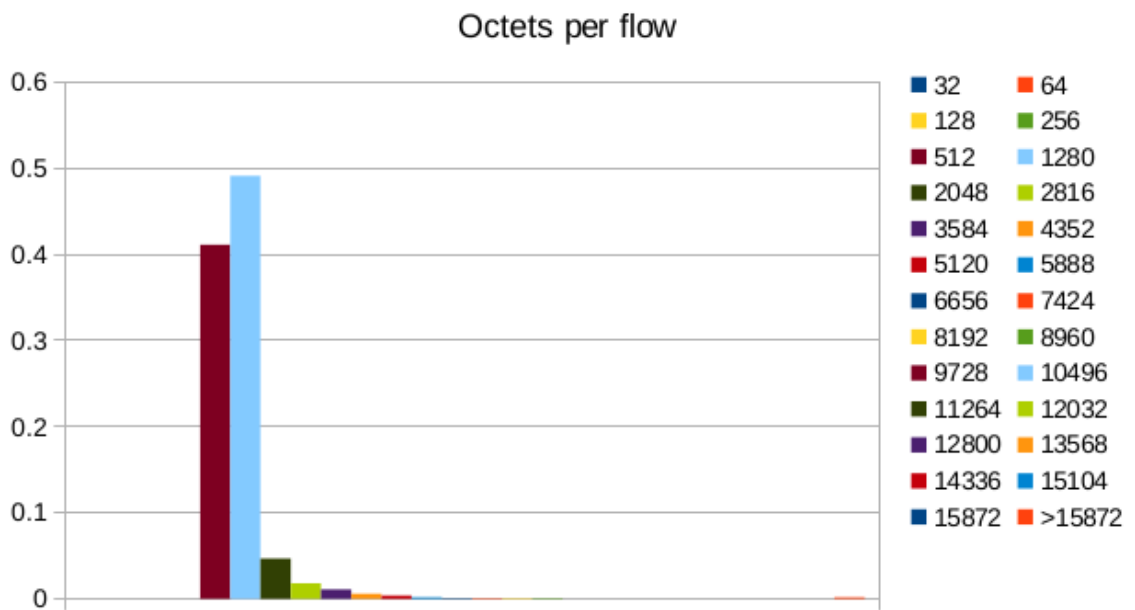
### 5b

```
root@group08-lg3:~# service snmpd start
root@group08-lg3:~# ps -eaf | grep snmpd
Debian+ 23773      1  0 08:32 ?                00:00:00 /usr/sbin/snmpd -Lsd -
          Lf /dev/null -u Debian-snmp -g Debian-snmp -I -smux mteTrigger
          mteTriggerConf -f
root      23834 20676   0 08:36 hvc0          00:00:00 grep snmpd

root@group08-lg3:~# snmpget -v1 -Of -c OUR_SECRET 10.8.3.1 iso.org.dod.
internet.mgmt.mib-2.system.sysUpTime.0
.iso.org.dod.internet.mgmt.mib-2.system.sysUpTime.sysUpTimeInstance =
Timeticks: (103560) 0:17:15.60
```

### 5c

UDP is used, because it behaves far better in networks, where a lot of packets are lost (UDP even begins to perform better at about 5%+ packet loss). In case of packet loss, the data is lost and cannot be recovered, because the sender does not know about the loss and does not keep a copy of






Figure 6: The packets per flow distribution with maximum throughput.

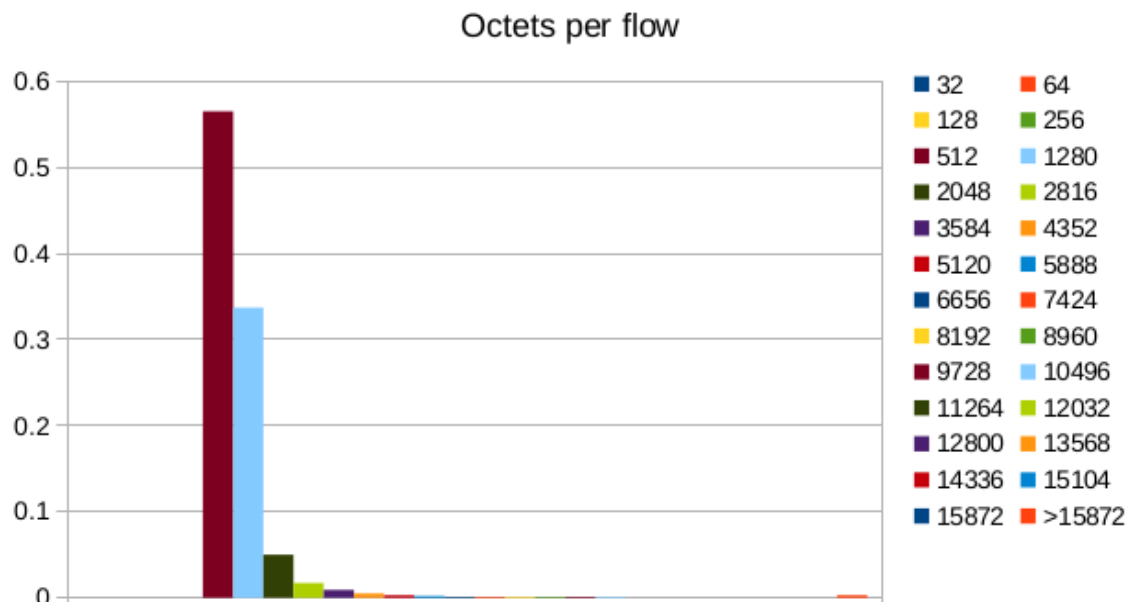


Figure 7: The octets per flow distribution with maximum throughput.

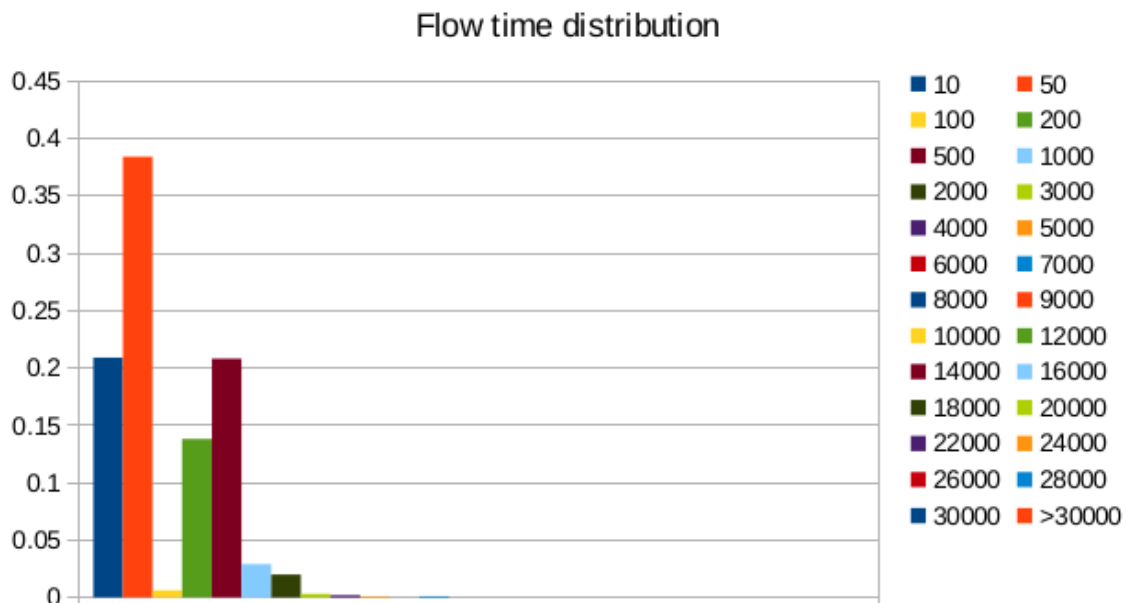


Figure 8: The flow time distribution with maximum throughput.

the packet. However, another request can be sent by the manager to get the current information (which might have changed in the mean time). In SNMP there also is the option to use SNMP Informs, which are acknowledged at application layer and therefore the most reliable choice if such acknowledgement is required.

If there is no agent running on the router, the sender will not receive an answer and it will appear as if the packet was lost. But for the sender it is impossible to tell the difference between packet loss, no running agent on the router and an incorrect community string.

The SNMP community string serves the purpose of authenticating the device requesting from or setting information on the agent. If it is incorrect the manager will not retrieve a notification that authentication failed but wait for a timeout:

```
root@group08-lg3:~# snmpget -v1 -Of -c WRONG_SECRET 10.8.3.1 iso.org.
dod.internet.mgmt.mib-2.system.sysUpTime.0
Timeout: No Response from 10.8.3.1.
```

Unfortunately, at least in version 1 all community strings are sent in cleartext, which makes SNMP an insecure protocol, especially if devices are not configured to be in read-only mode (as we did).

## 5d

We use the following command to capture the packets:

```
root@group08-lg3:~# tshark -i eth1 -w capture.pcap
Running as user "root" and group "root". This could be dangerous.
tshark: Lua: Error during loading:
[string "/usr/share/wireshark/init.lua"]:44: dofile has been disabled
due to running Wireshark as superuser. See https://wiki.wireshark.
org/CaptureSetup/CapturePrivileges for help in running Wireshark
as an unprivileged user.
Capturing on 'eth1'
[2702983.326642] device eth1 entered promiscuous mode
[...]
```

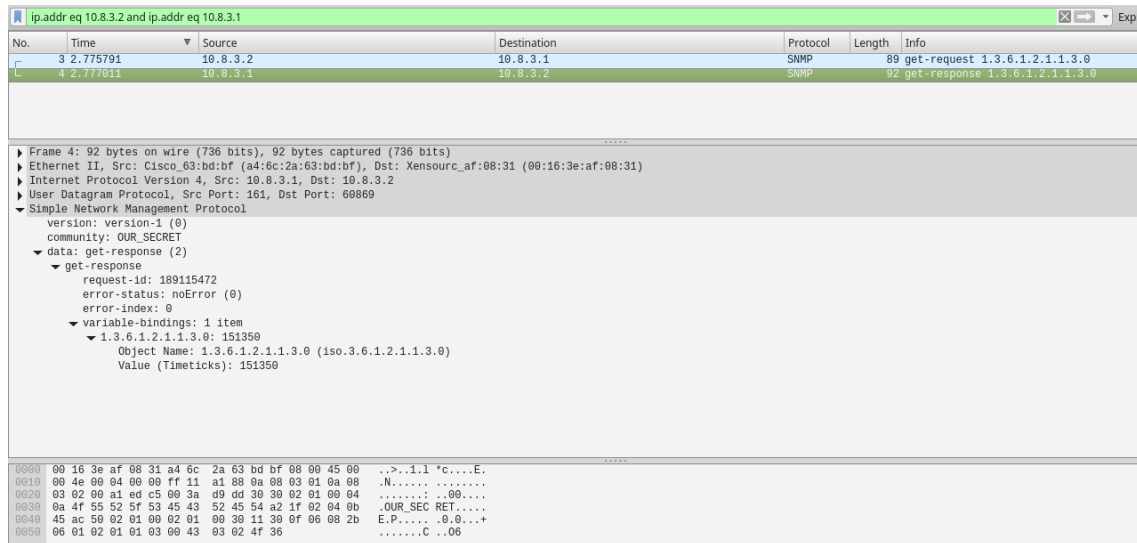


Figure 9: The SNMPget response.

Then we create an `snmpget` request as shown above and capture it. We analyze everything in Wireshark using a filter to get rid of non-SNMP packets (see Figure 9). Obviously we see the usual protocol stack with UDP at the transport layer. On top of it we find SNMP. We can see the community secret and the version in plaintext. Wireshark shows the message type and the data payload. The main difference between request and response is the value, which is displayed in the response. It is the number of timeticks. On the command line SNMP displays this value both as the number of time ticks and as a human readable time, along with the OID where it was found. In the packet the OID is a numeric value, while it is a human readable string in the cli. It is also possible to provide the numerical value, when requesting the OID.

## 5e

We set the MIBs on `rc1`:

```
lev-rc1(config)#snmp-server location sweethome
lev-rc1(config)#snmp-server contact ourcontact
```

Then we do the `snmpwalk` and capture it with `tshark` as above. The full output is provided in an extra file in the attachments as it is more than 800 lines:

```
root@group08-1g3:~# snmpwalk -v1 -Of -c OUR_SECRET 10.8.3.1 iso.org.dod
.internet.mgmt.mib-2.system
[...]
.iso.org.dod.internet.mgmt.mib-2.system.sysUpTime.sysUpTimeInstance =
  Timeticks: (414864) 1:09:08.64
.iso.org.dod.internet.mgmt.mib-2.system.sysContact.0 = STRING:
  ourcontact
.iso.org.dod.internet.mgmt.mib-2.system.sysName.0 = STRING: lev-rc1
.iso.org.dod.internet.mgmt.mib-2.system.sysLocation.0 = STRING:
  sweethome
.iso.org.dod.internet.mgmt.mib-2.system.sysServices.0 = INTEGER: 78
[...]
```

`snmpwalk` basically saves users time by always automatically issuing the `get-next-request`, as it is shown by Wireshark (see attached capture file and Figure 10). Therefore it is used to retrieve a complete subtree of information with just one call instead of calling each unit of information with an individual `get` request. It will stop when the information requested is no longer part of the requested OID. If it is performed on the system OID it usually retrieves the complete tree of

No.	Time	Source	Destination	Protocol	Length	Info
5	5.012599921	10.8.3.1	10.8.3.2	SNMP	343	get-response 1.3.6.1.2.1.1.1.0
6	5.012791721	10.8.3.2	10.8.3.1	SNMP	89	get-next-request 1.3.6.1.2.1.1.1.0
7	5.013762809	10.8.3.1	10.8.3.2	SNMP	98	get-response 1.3.6.1.2.1.1.2.0
8	5.013871813	10.8.3.2	10.8.3.1	SNMP	89	get-next-request 1.3.6.1.2.1.1.2.0
9	5.014776295	10.8.3.1	10.8.3.2	SNMP	92	get-response 1.3.6.1.2.1.1.3.0
10	5.014861871	10.8.3.2	10.8.3.1	SNMP	89	get-next-request 1.3.6.1.2.1.1.3.0
11	5.015708837	10.8.3.1	10.8.3.2	SNMP	99	get-response 1.3.6.1.2.1.1.4.0
12	5.015787142	10.8.3.2	10.8.3.1	SNMP	89	get-next-request 1.3.6.1.2.1.1.4.0
13	5.016643501	10.8.3.1	10.8.3.2	SNMP	96	get-response 1.3.6.1.2.1.1.5.0
14	5.016724583	10.8.3.2	10.8.3.1	SNMP	89	get-next-request 1.3.6.1.2.1.1.5.0
15	5.017566813	10.8.3.1	10.8.3.2	SNMP	98	get-response 1.3.6.1.2.1.1.6.0
16	5.017642905	10.8.3.2	10.8.3.1	SNMP	89	get-next-request 1.3.6.1.2.1.1.6.0

▶ Frame 15: 98 bytes on wire (784 bits), 98 bytes captured (784 bits) on interface 0  
 ▶ Ethernet II, Src: Cisco\_63:bd:bf (a4:6c:2a:63:bd:bf), Dst: Xensourc\_af:08:31 (00:16:3e:af:08:31)  
 ▶ Internet Protocol Version 4, Src: 10.8.3.1, Dst: 10.8.3.2  
 ▶ User Datagram Protocol, Src Port: 161, Dst Port: 53962  
 ▶ Simple Network Management Protocol  
   version: version-1 (0)  
   community: OUR\_SECRET  
   data: get-response (2)  
     get-response  
       request-id: 1155176804  
       error-status: noError (0)  
       error-index: 0  
       variable-bindings: 1 item  
         1.3.6.1.2.1.1.6.0: 7377656574686fd65  
           Object Name: 1.3.6.1.2.1.1.6.0 (iso.3.6.1.2.1.1.6.0)  
           Value (OctetString): 7377656574686fd65

0000 00 16 3e af 08 31 a4 6c 2a 63 bd bf 08 00 45 00 ...1.1 "c....E.  
 0010 00 54 01 5a 00 00 ff 11 a0 2c 0a 08 03 01 0a 08 .T.Z... ..  
 0020 03 02 00 a1 d2 ca 00 00 cf f2 30 36 02 01 00 04 .....@ ..06...  
 0030 0a 4f 55 52 5f 53 45 43 52 45 54 a2 25 02 04 44 .OUR\_SEC RET%.D  
 0040 0a 99 64 02 01 00 02 01 00 30 17 30 15 06 00 20 .....00...+  
 0050 06 01 02 01 01 06 00 04 09 73 77 65 65 74 68 6f d6 5 .7377656574686fd65  
 0060 73 77 65 65 74 68 6f d6 5

Figure 10: The SNMPwalk capture in Wireshark.

information stored on the machine in the system MIB group. AS we can see in the capture each snmp-next-request is respondede with the same kind of snmp-response as in 5d.

## Included files

capture\_question\_05d.pcap, capture\_question\_05e.pcap, snmpwalk.txt, q03-config-rc1.txt, q03-config-sc1.txt, q04-config-rc1.txt